

Assignment Project Exam Help

Add WeChat powcoder

24. Virtual Memory: TLB and Caches

Assignment Project Exam Help

EECS 370 – Introduction to Computer Organization – Fall 2020

<https://powcoder.com>

Satish Narayanasamy
Add WeChat powcoder

EECS Department
University of Michigan in Ann Arbor, USA

© Narayanasamy 2020

The material in this presentation cannot be
copied in any form without written permission

Assignment Project Exam Help

Final Exam

Add WeChat powcoder

Online exam through Gradescope

Practice exam on Gradescope will be made available

Assignment Project Exam Help

Topics

Strong emphasis on topics since the midterm

<https://powcoder.com>

Pipelining

Branch prediction

Caches

Virtual memory

Add WeChat powcoder

Assignment Project Exam Help


$$\text{Page offset size} = \log(4 \text{ KB}) = 12 \text{ bits}$$

Page size = 4 KB

Virtual Memory: An Example



Virtual memory
(2^{20} bytes = 256 pages)

Virtual address




Physical Memory
(16 KB = 4 pages)

Physical address



Virtual memory: 2^{20} bytes

$$2^{20} / 4 \text{ KB} = 256 \text{ pages}$$




Page Table
(256 entries)

Assignment Project Exam Help


<https://powcoder.com>

Add WeChat powcoder



Pinned

Physical Memory: 16 KB
(16 KB / 4 KB = 4 pages)



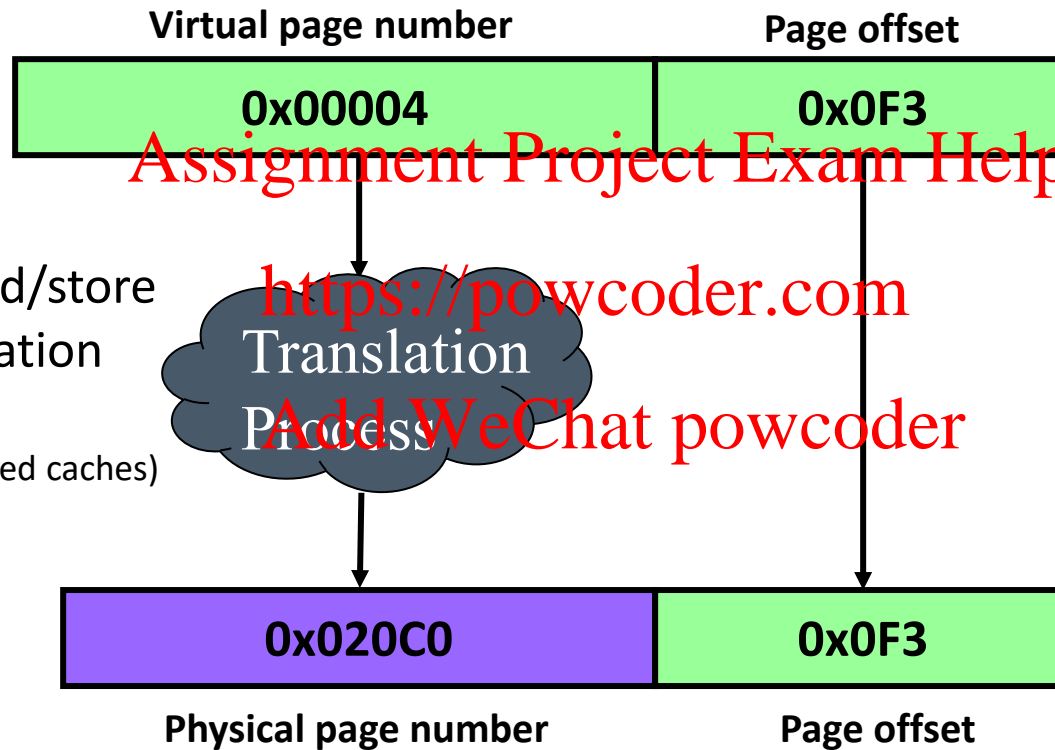
Disk
(swap partition)

Assignment Project Exam Help

Address Translation

Add WeChat powcoder

Virtual address = 0x000040F3



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Every instruction fetch, load/store
needs to do address translation

(optimized later with virtually-addressed caches)

Physical address = 0x020C00F3

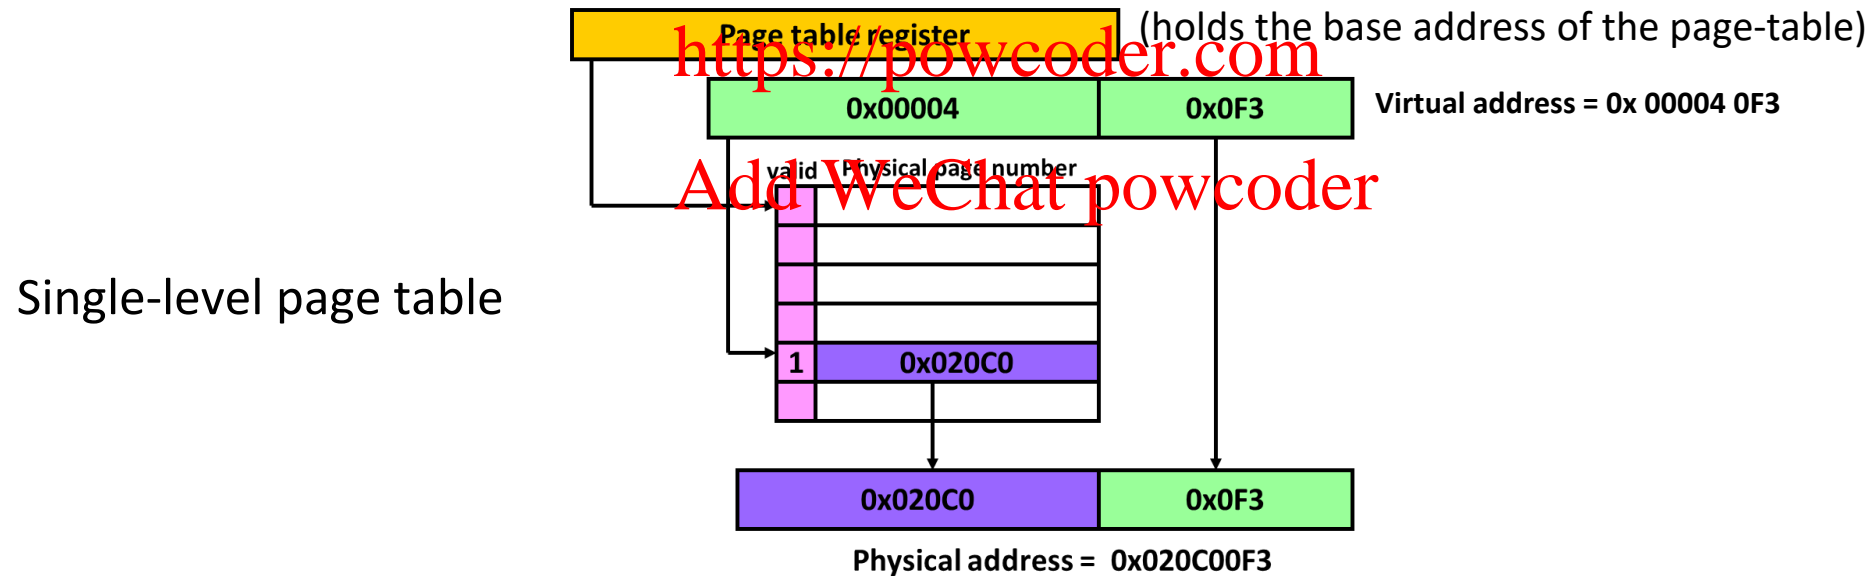
Assignment Project Exam Help

Page table lookups for address translation

N-level page table

Each address translation requires **N memory accesses**, one per level

Assignment Project Exam Help



Assignment Project Exam Help

Problem: Address translation overhead

Address translation is on the critical path

Can be done only after virtual address is known

Need to done before accessing memory (optimized later with virtually-addressed caches)

Assignment Project Exam Help

Address translation requires accesses to the page table(s) in physical memory

A memory access (instruction fetch, load/store) performs **N additional memory accesses**,

where N is the number of levels in a hierarchical page-table

Slow

After address translation, memory hierarchy is accessed to perform memory access

Solution: Translation look-aside buffer (TLB)

Add WeChat powcoder

TLB is a special cache for page-tables.

Speeds-up address translation by reducing main memory accesses to page tables.

On a TLB miss, access page-table(s) in main memory

Assignment Project Exam Help

<https://powcoder.com>

Stores a small subset of valid page table entries.

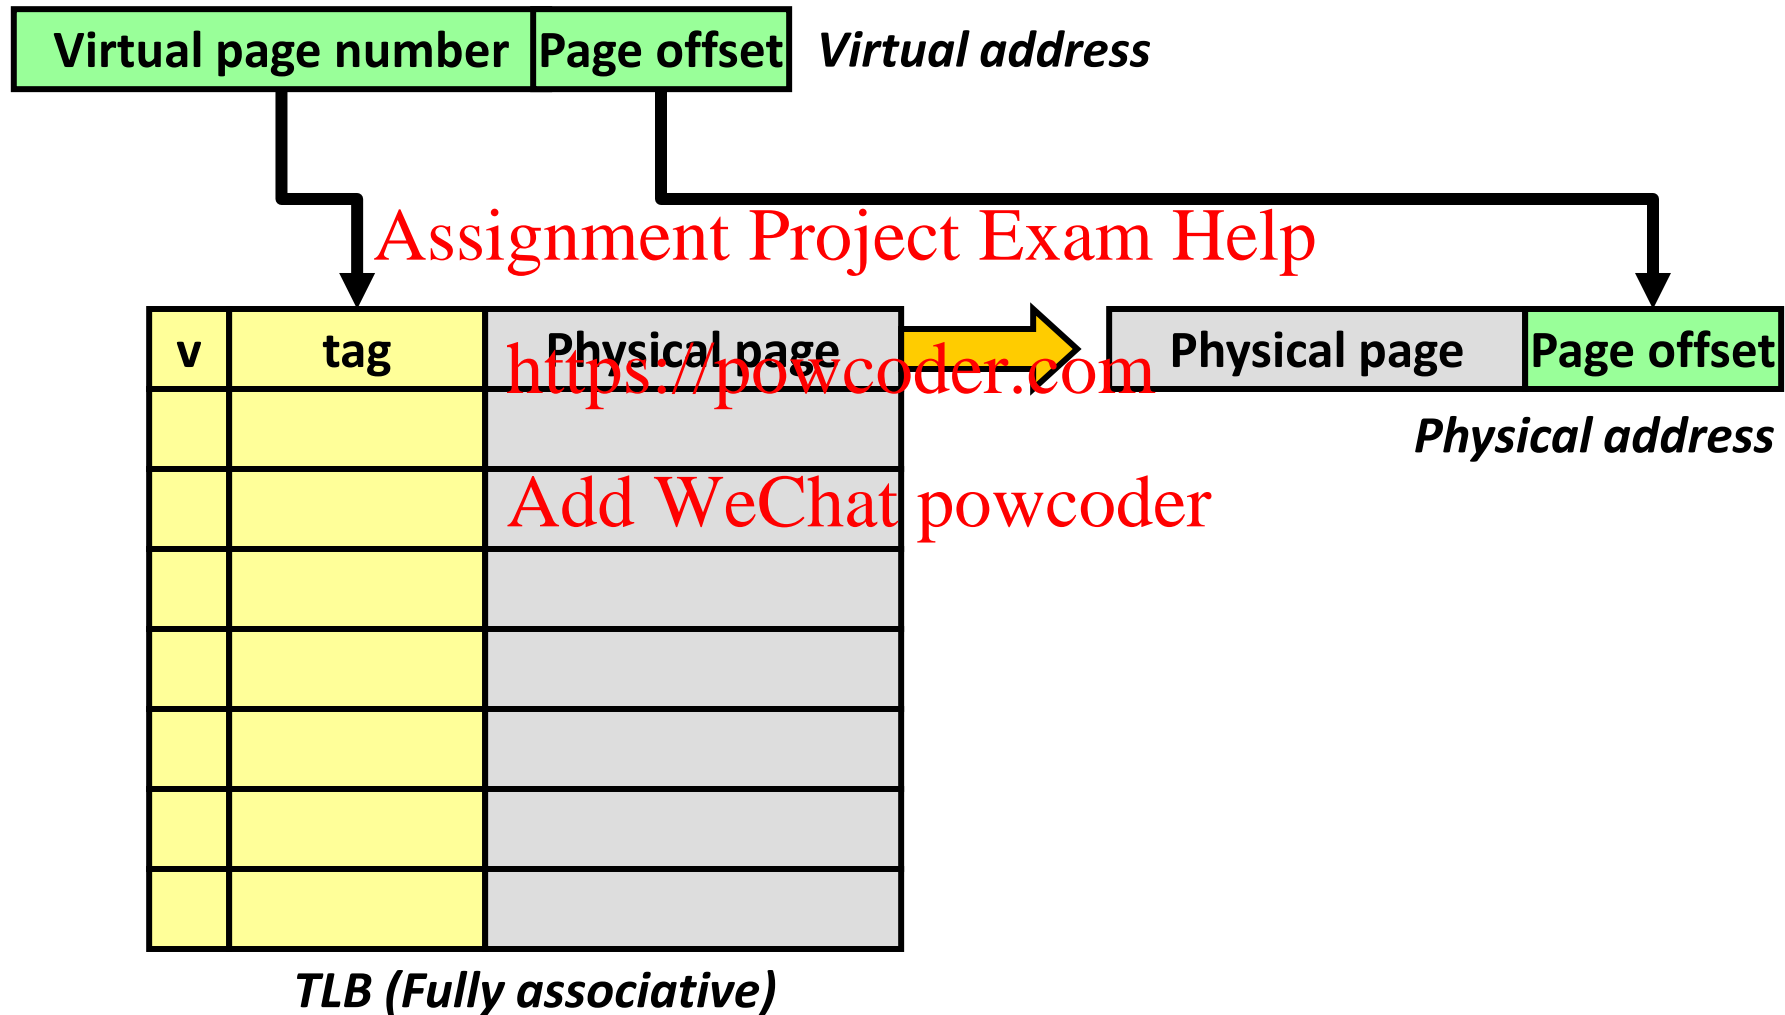
Add WeChat powcoder

16-512 entries common.

Typically, has low miss rate ($< 1\%$).

Assignment Project Exam Help Translation look-aside buffer (TLB)

Add WeChat powcoder



Assignment Project Exam Help

Putting it all together

Add WeChat powcoder

OS: loading program in memory

Creates a new process P

Constructs a page table for P

Assignment Project Exam Help

Marks all page table entries as invalid with a pointer to the disk image of the program

<https://powcoder.com>

That is, point to the executable file containing the binary.

Add WeChat powcoder

Runs the program

Will get an immediate page fault on the first instruction (everything is on disk initially)

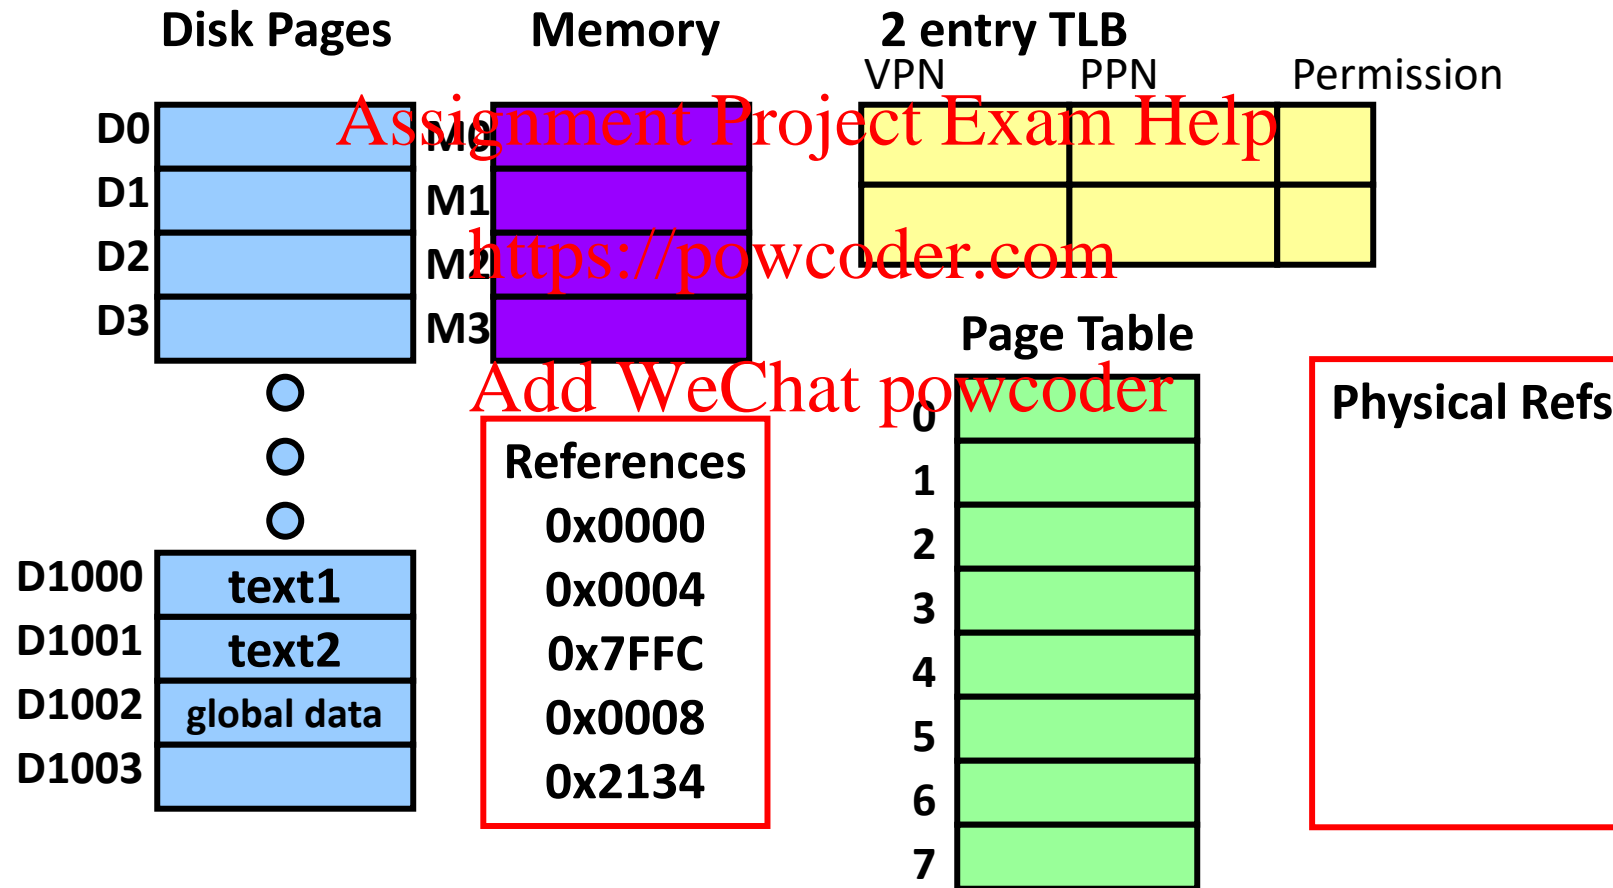
Assignment Project Exam Help

Loading a program into memory

Add WeChat powcoder

Page size = 4 KB, Page table entry size = 4 B

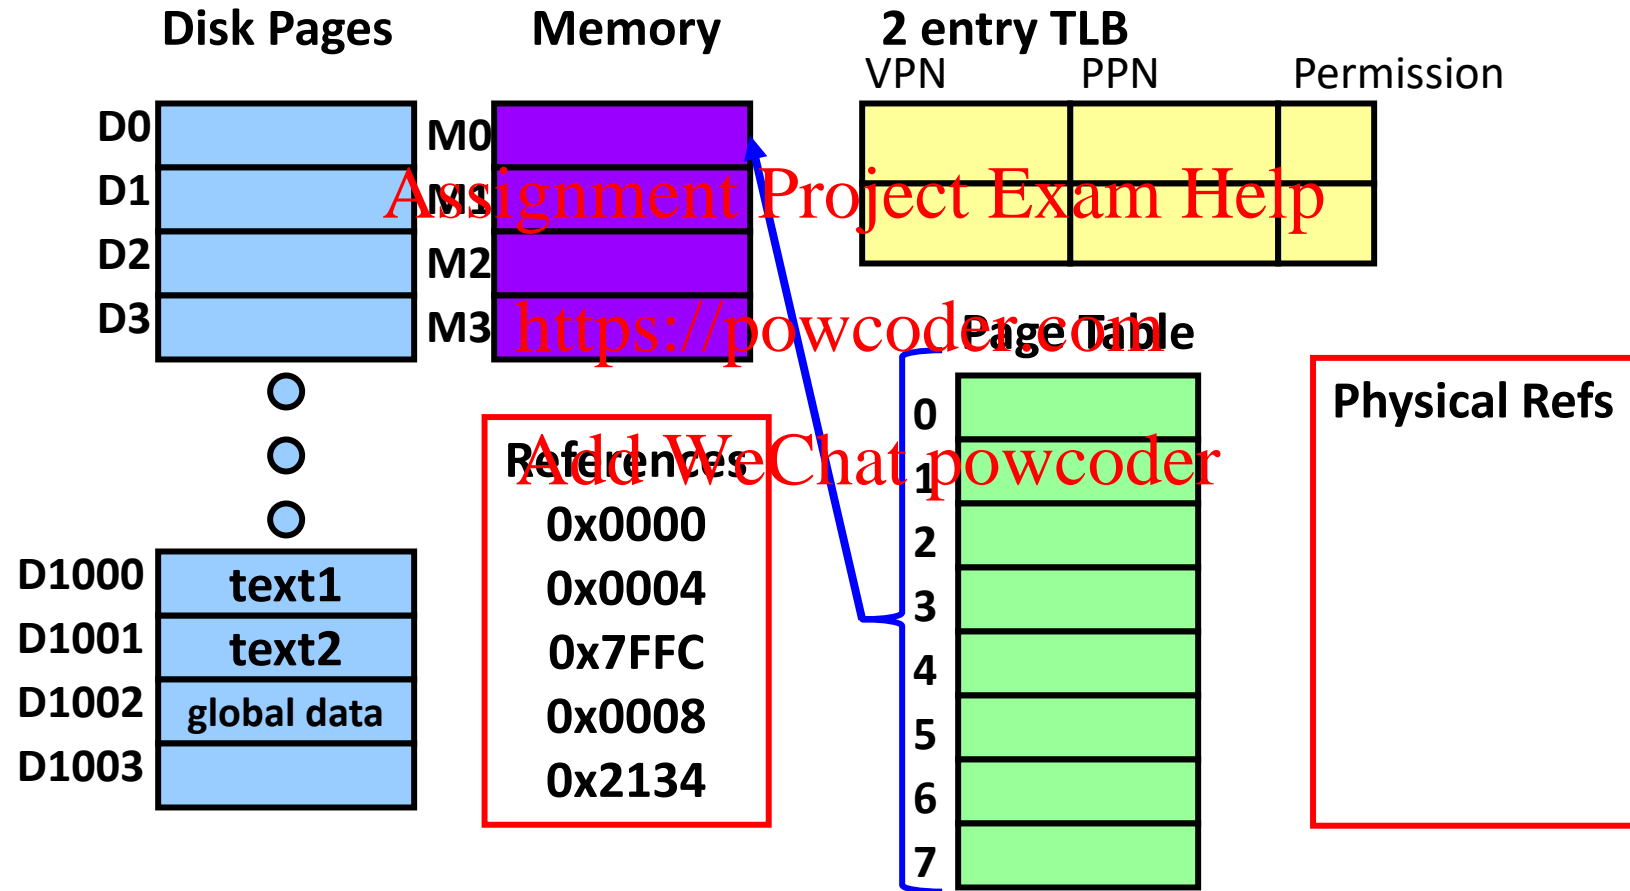
Page table register points to physical address 0x0000



Assignment Project Exam Help

Loading a program into memory

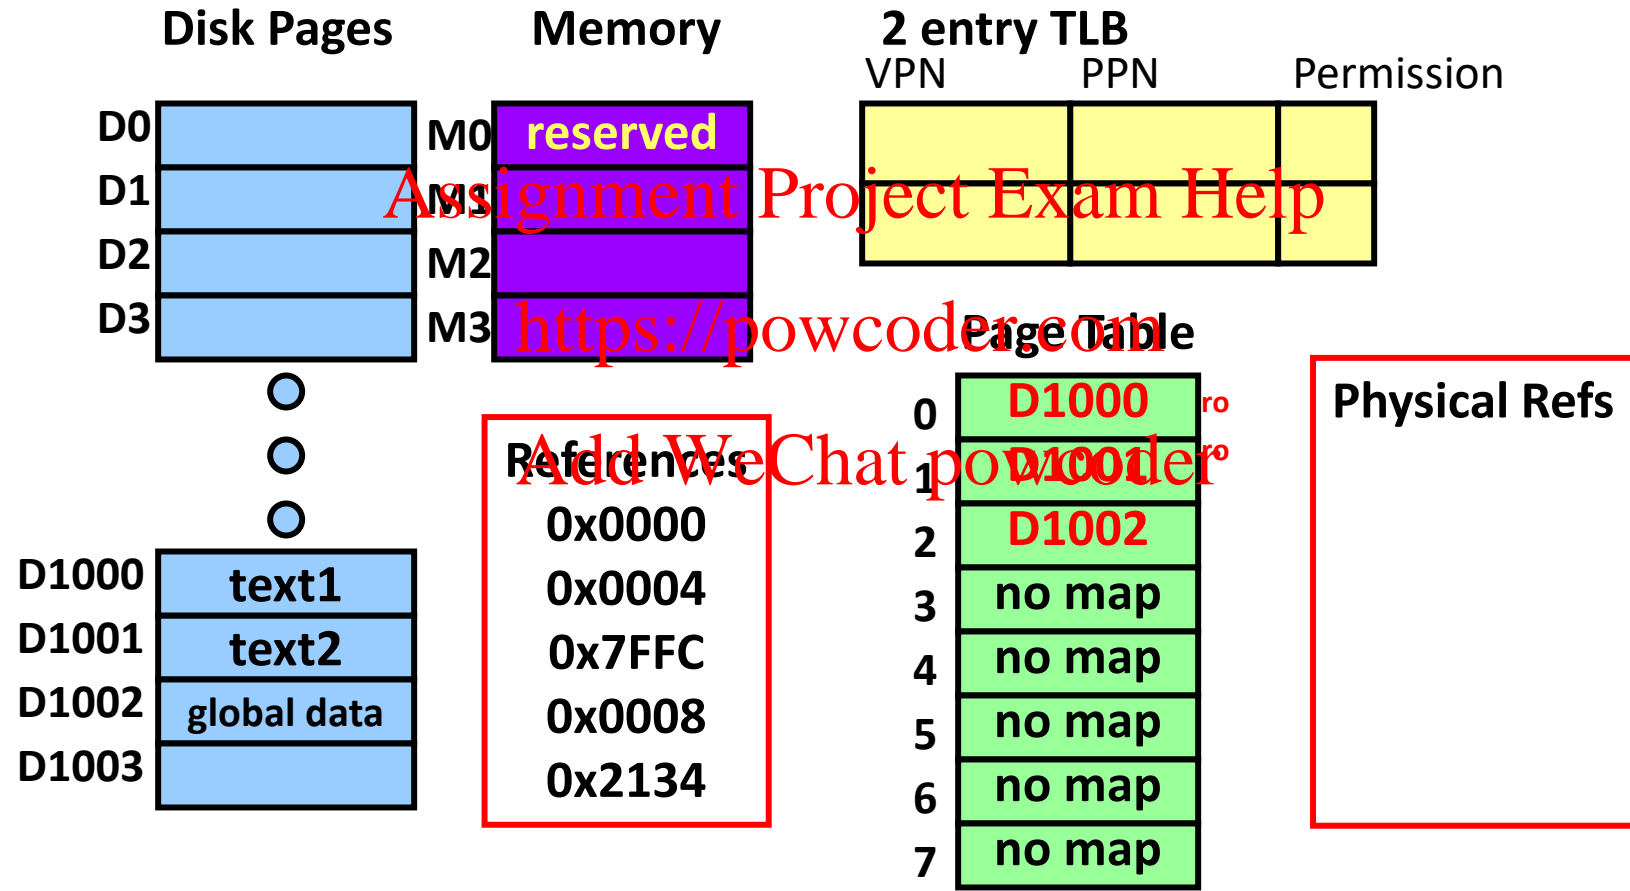
Add WeChat powcoder



Assignment Project Exam Help

Step 1: Read executable header & initialize page table

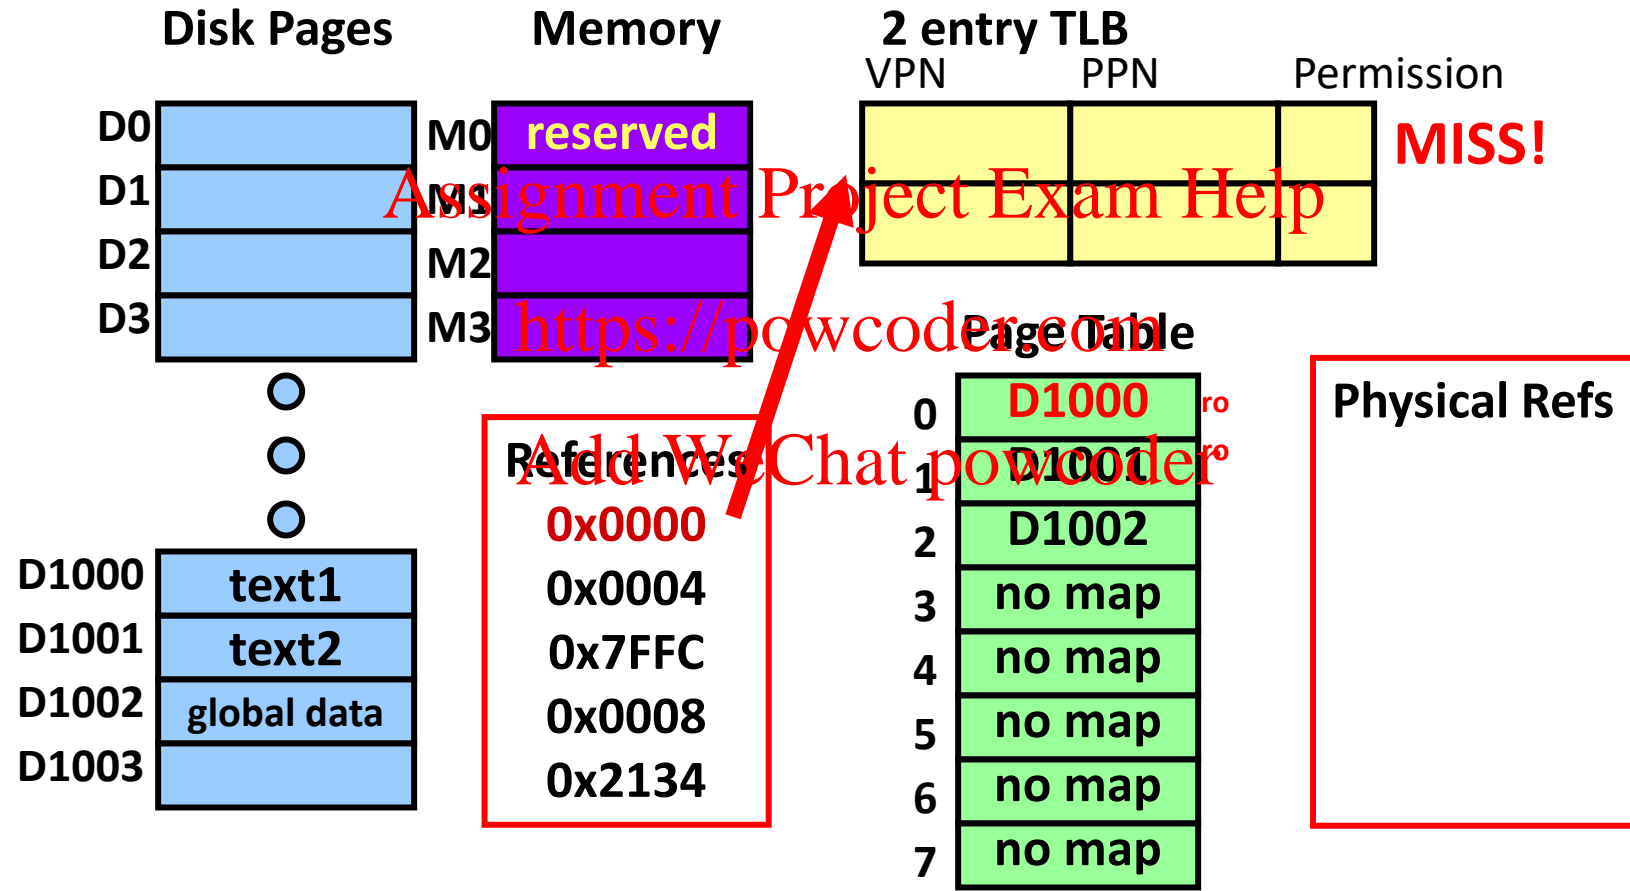
Add WeChat powcoder



Assignment Project Exam Help

Step 2: Load PC from header & start execution

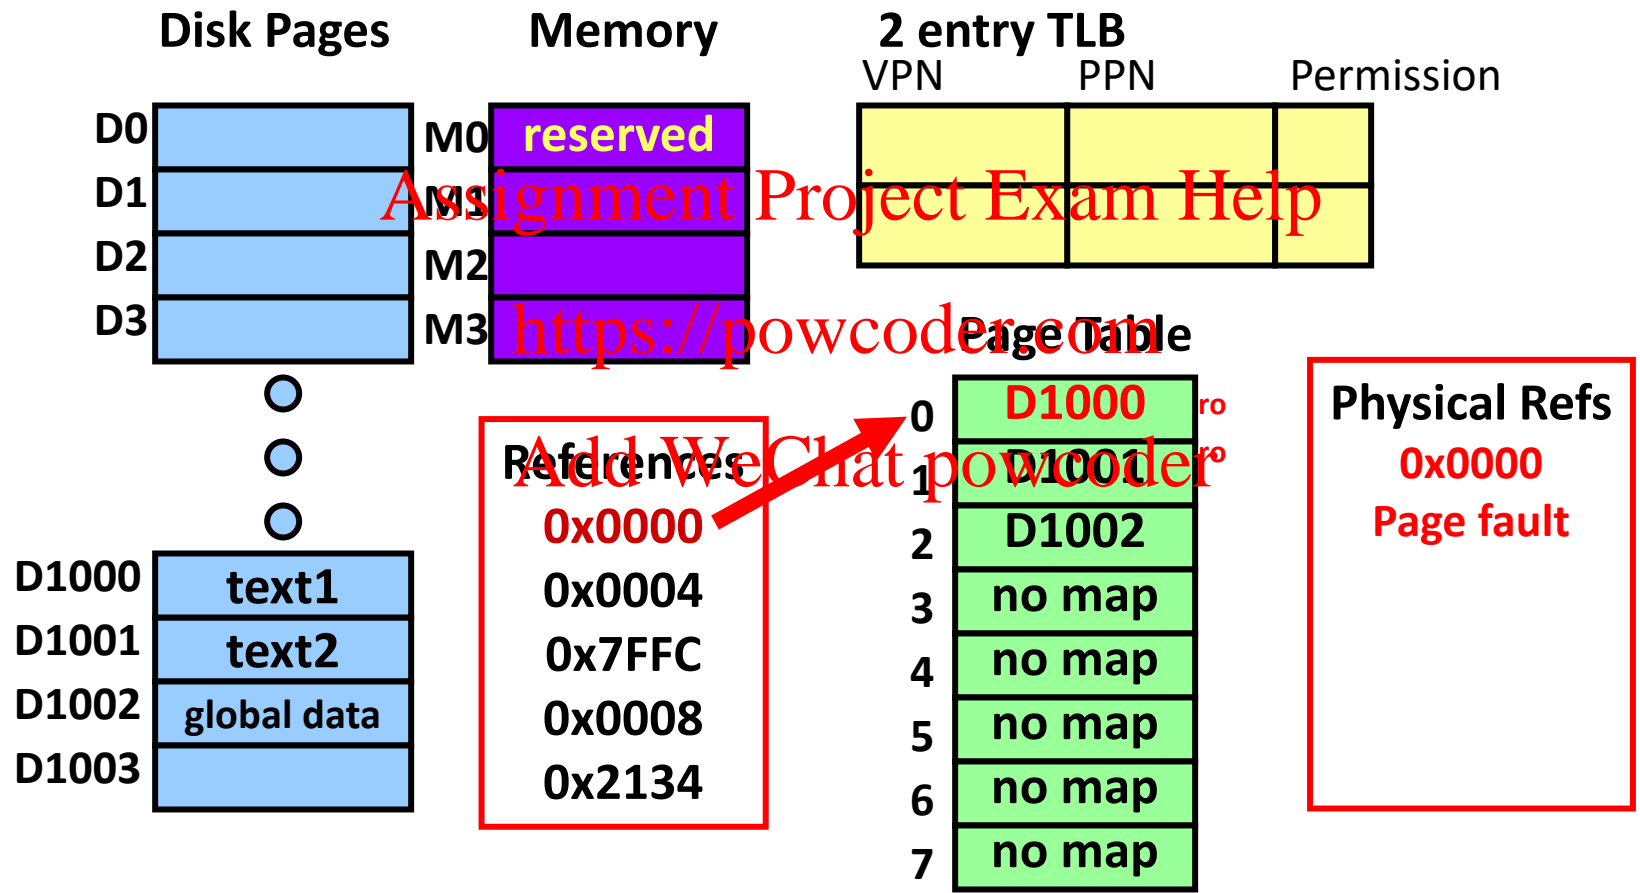
Add WeChat powcoder



Assignment Project Exam Help

Fetching instruction 0000

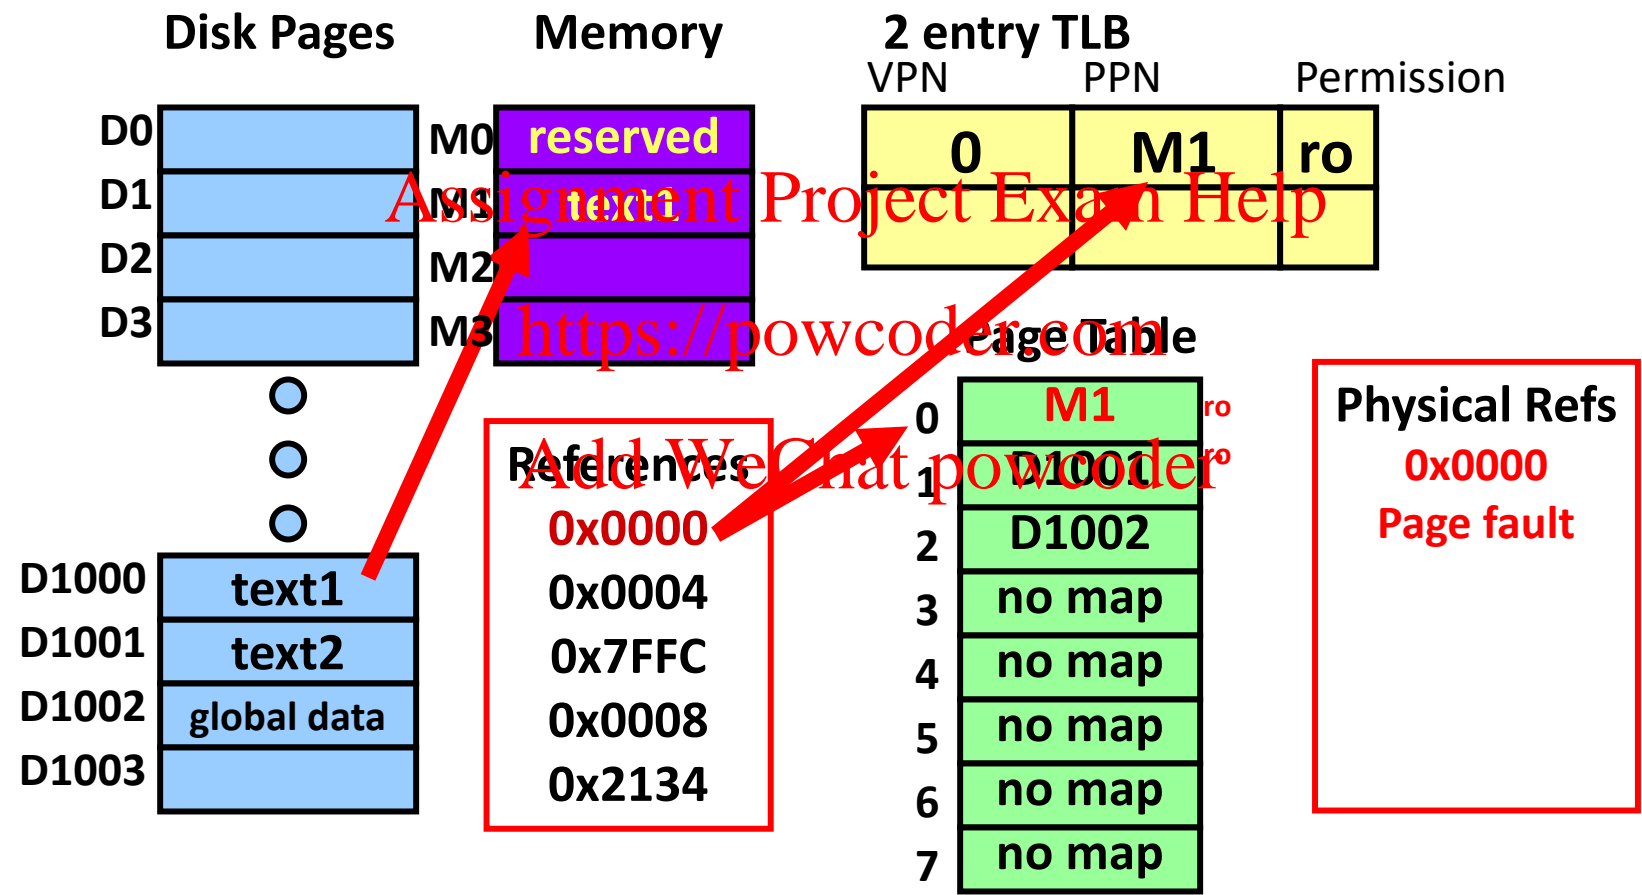
Add WeChat powcoder



Assignment Project Exam Help

Fetching instruction 0000

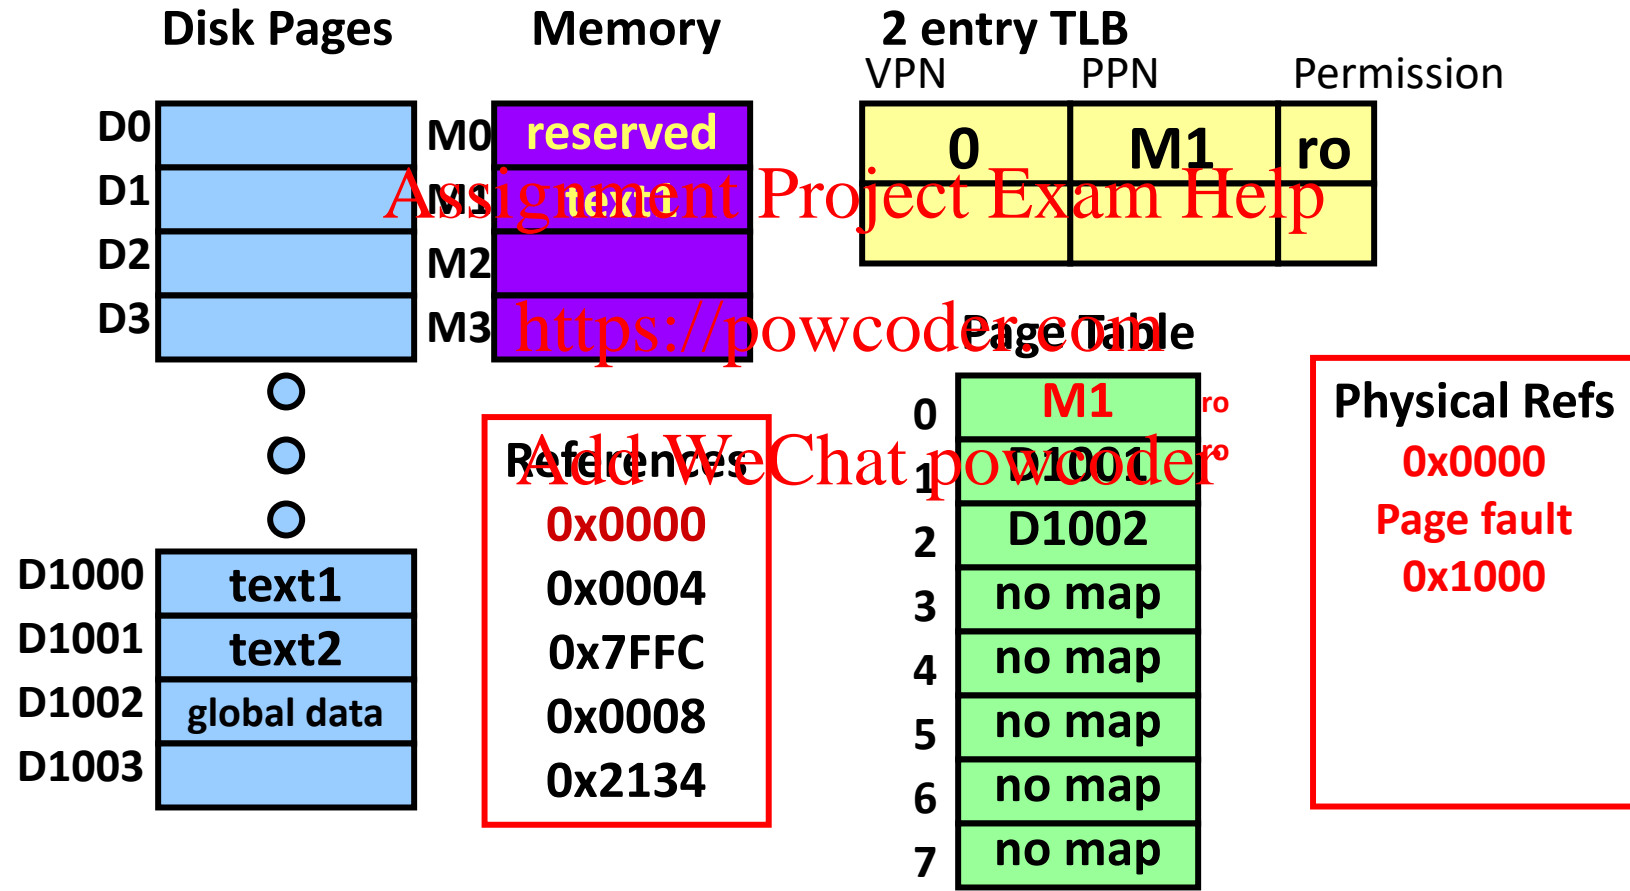
Add WeChat powcoder



Assignment Project Exam Help

Fetching instruction 0000

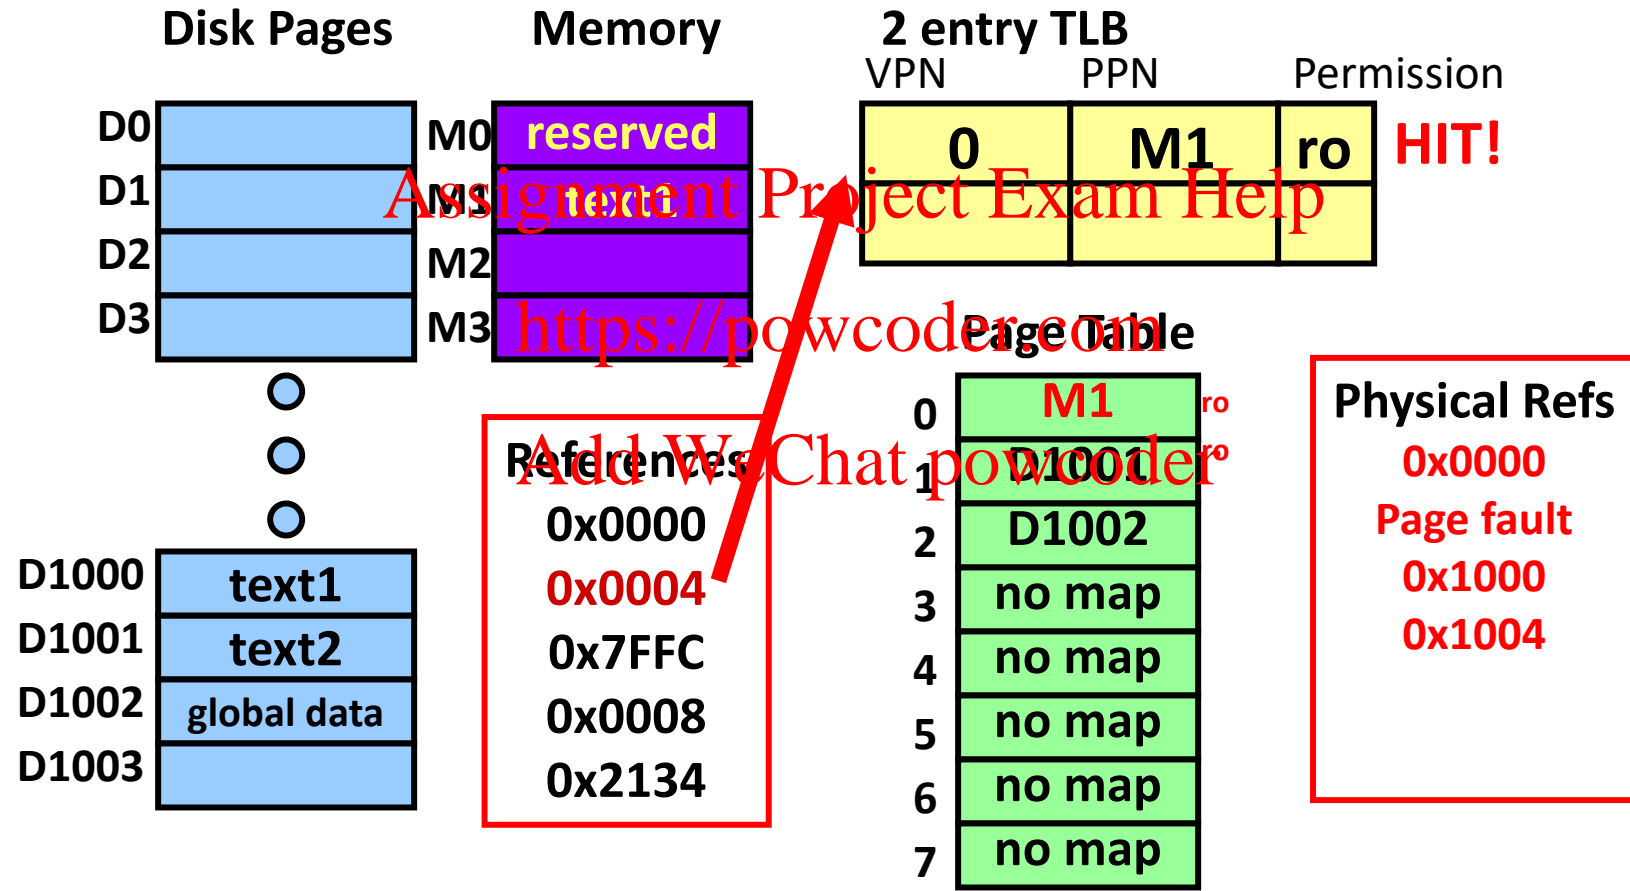
Add WeChat powcoder



Assignment Project Exam Help

Fetching instruction 0004

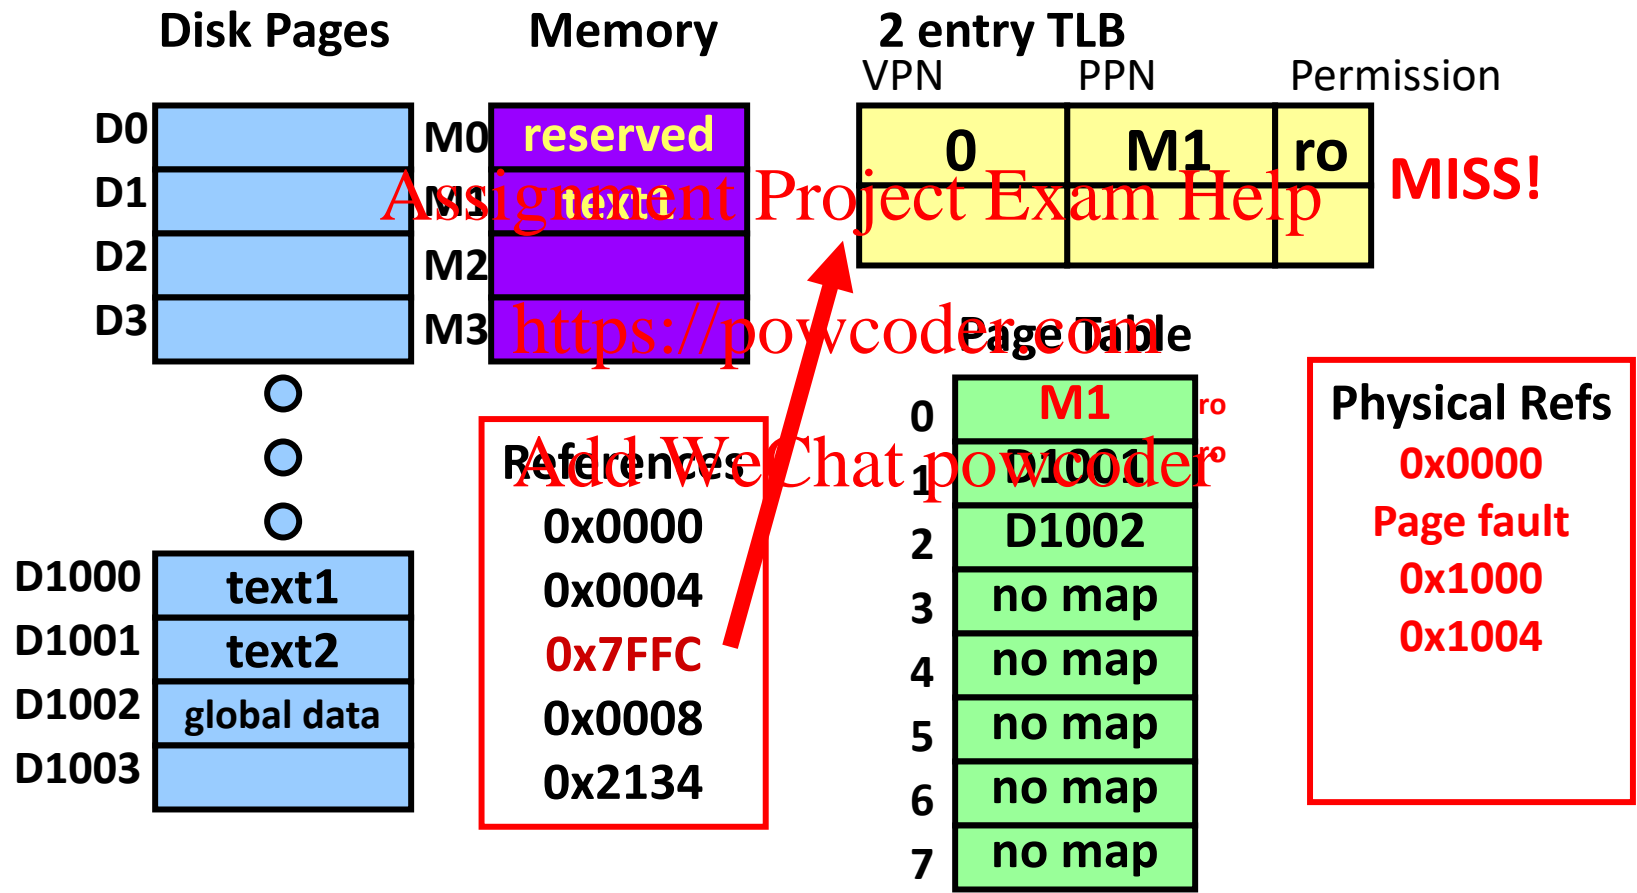
Add WeChat powcoder



Assignment Project Exam Help

Reference 7FFC

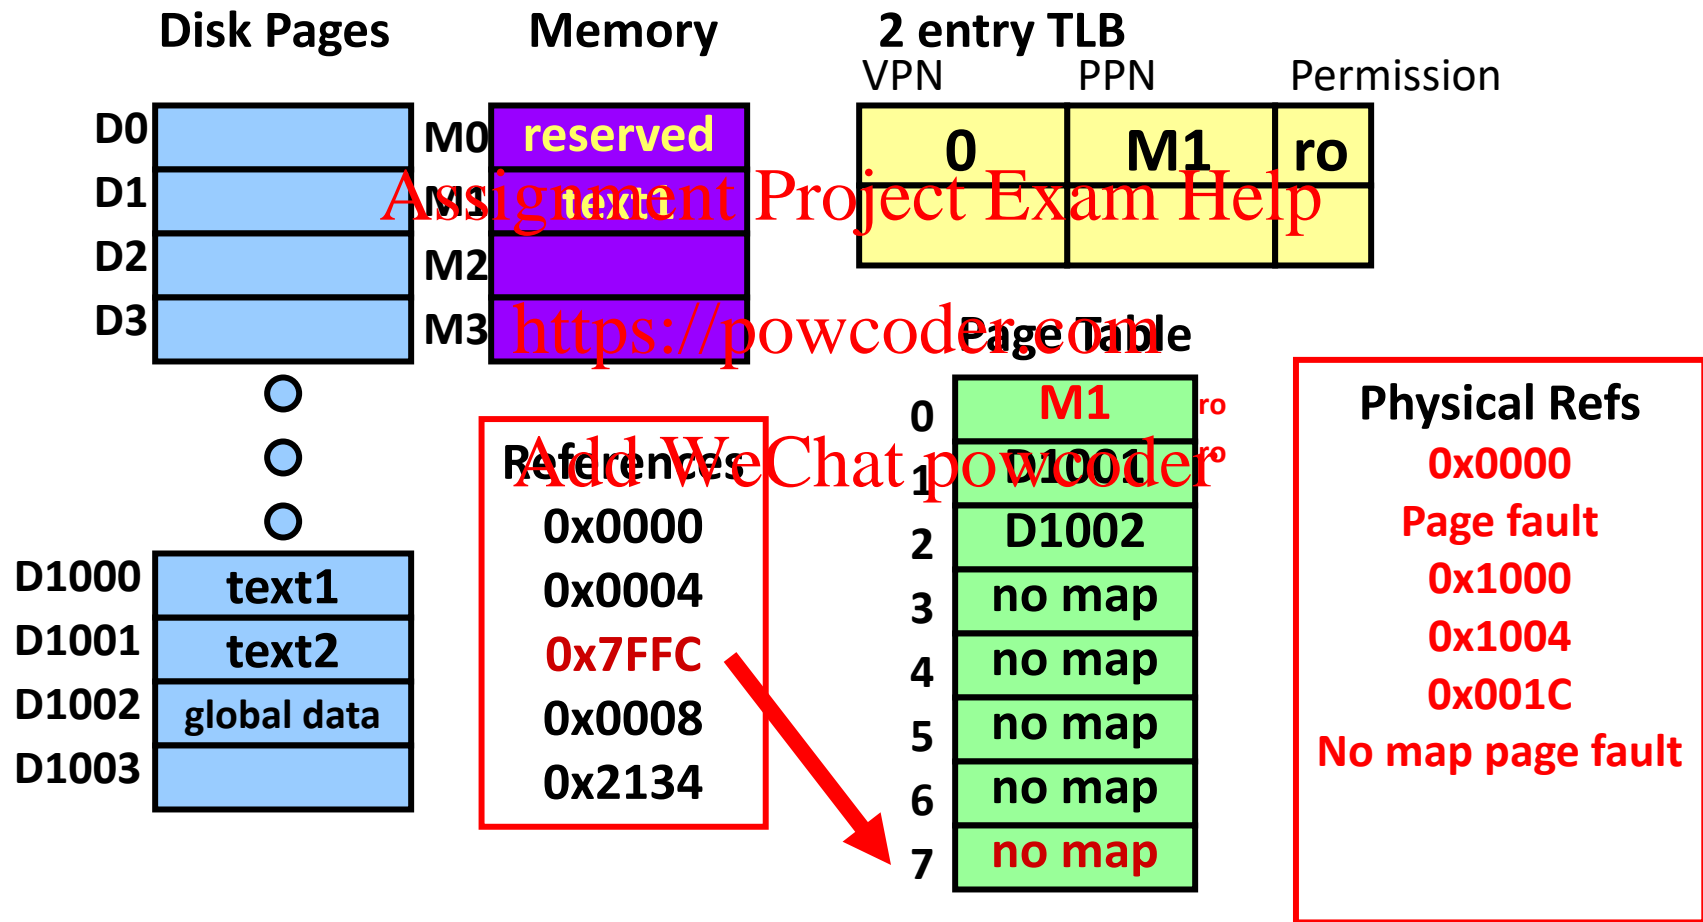
Add WeChat powcoder



Assignment Project Exam Help

Reference 7FFC

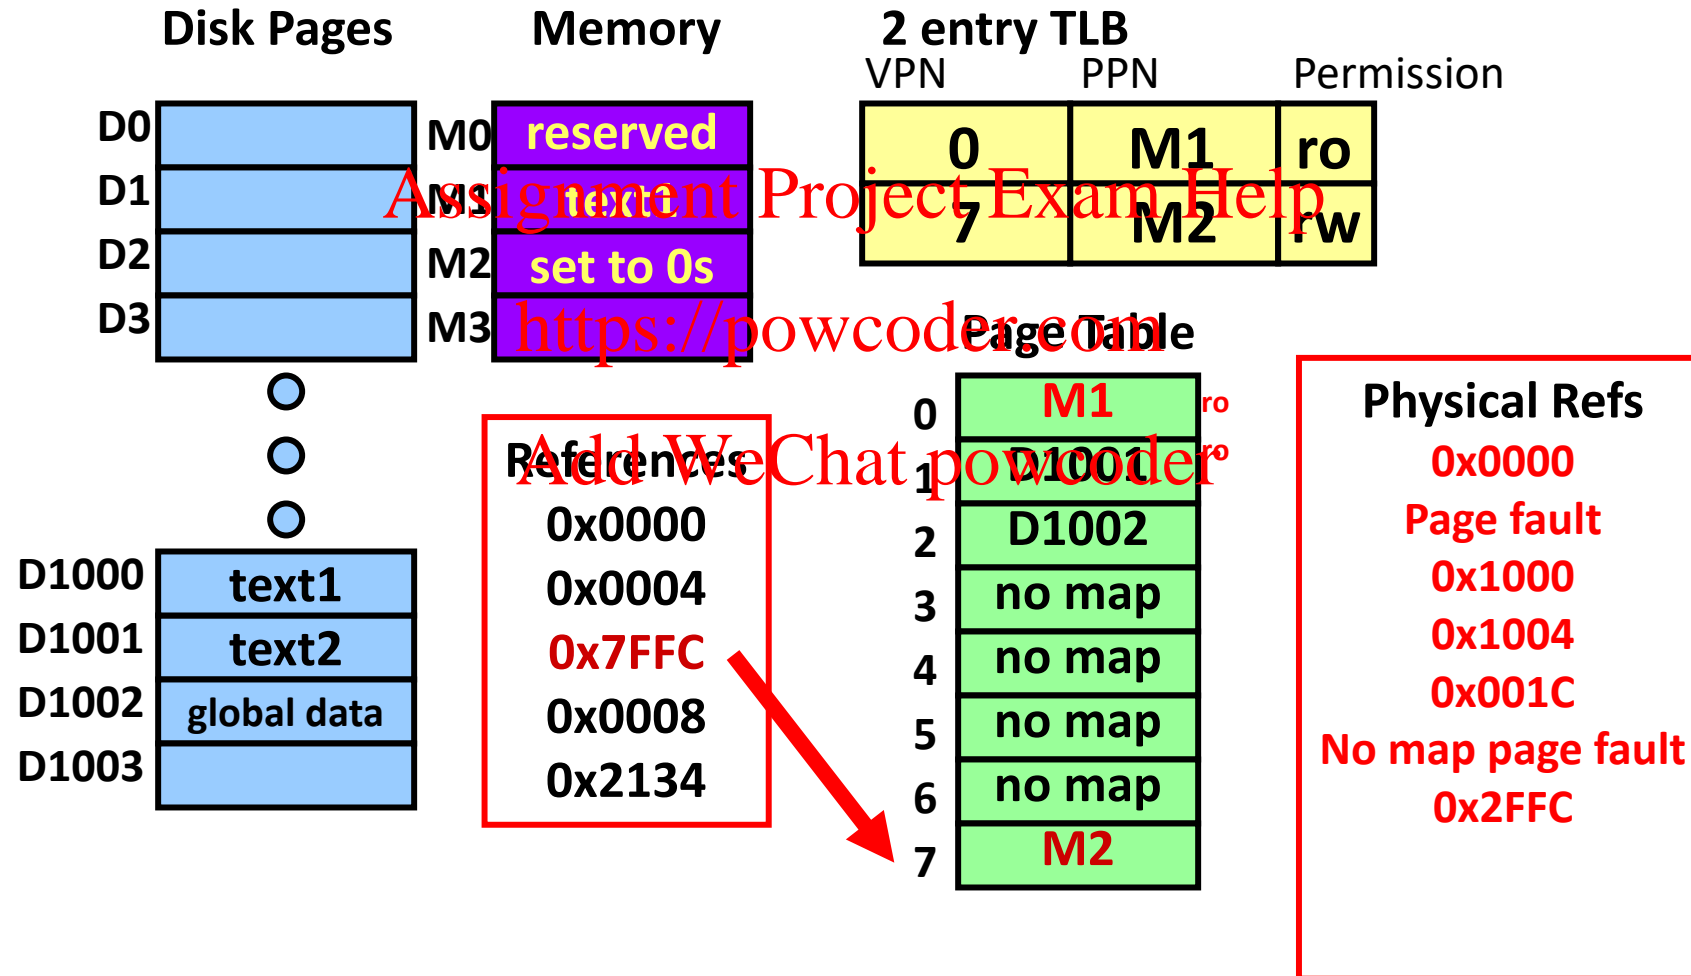
Add WeChat powcoder



Assignment Project Exam Help

Reference 7FFC

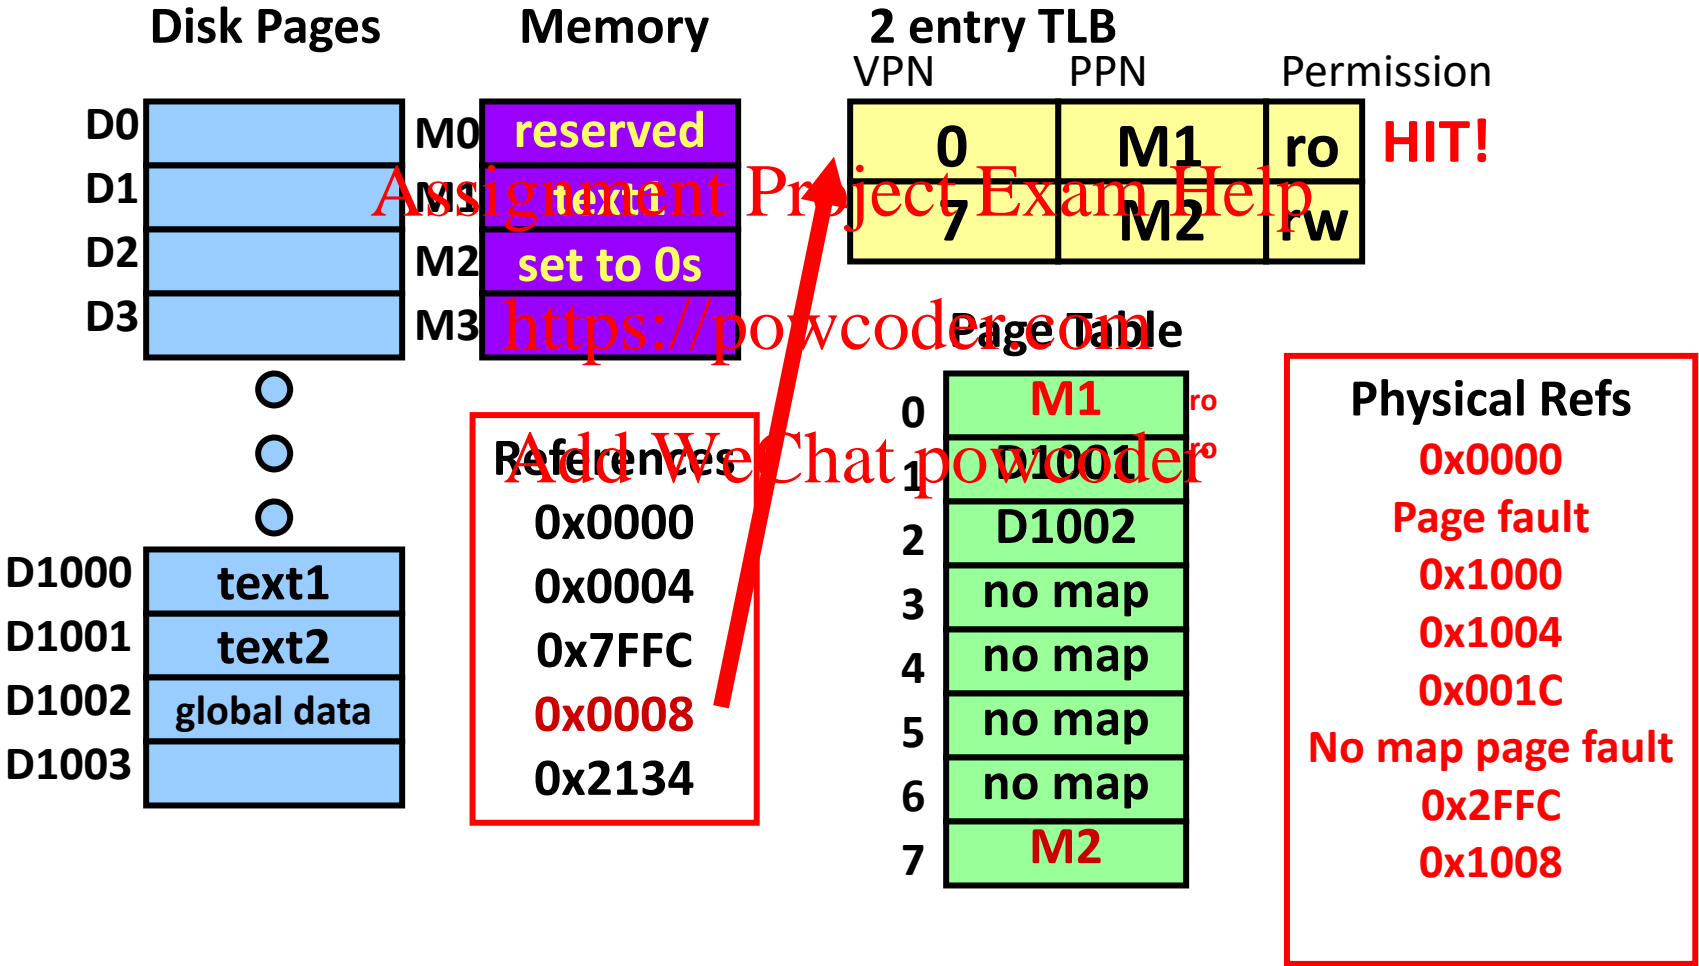
Add WeChat powcoder



Assignment Project Exam Help

Fetching instruction 0008

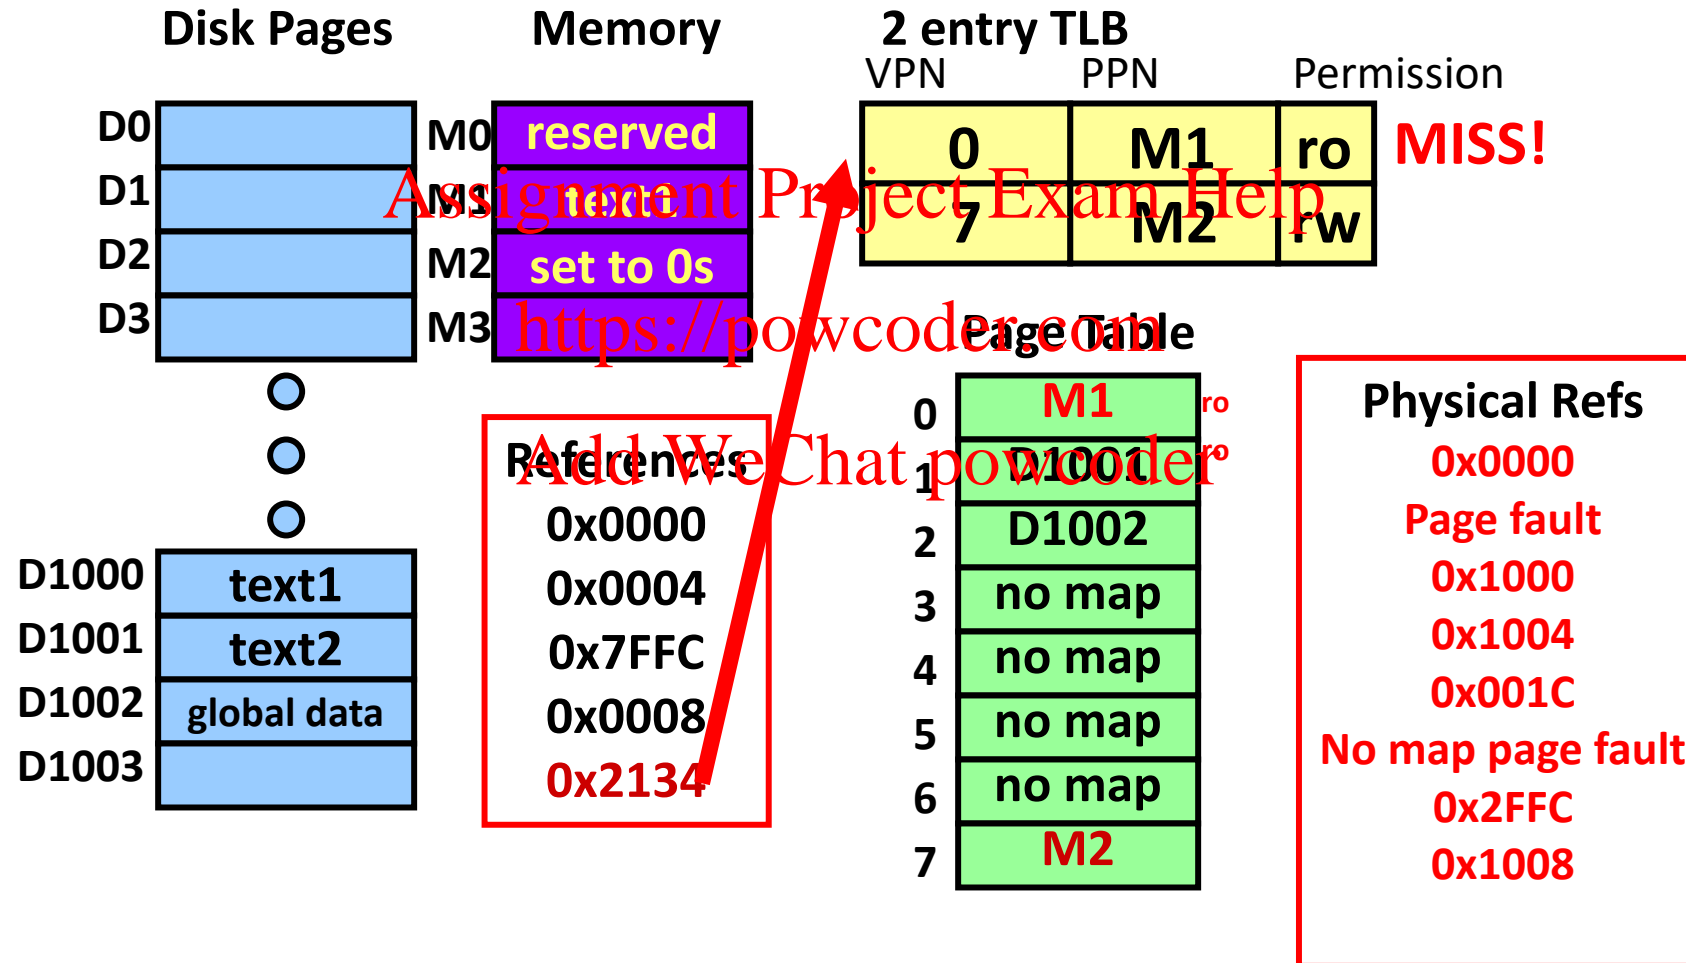
Add WeChat powcoder



Assignment Project Exam Help

Reference 2134

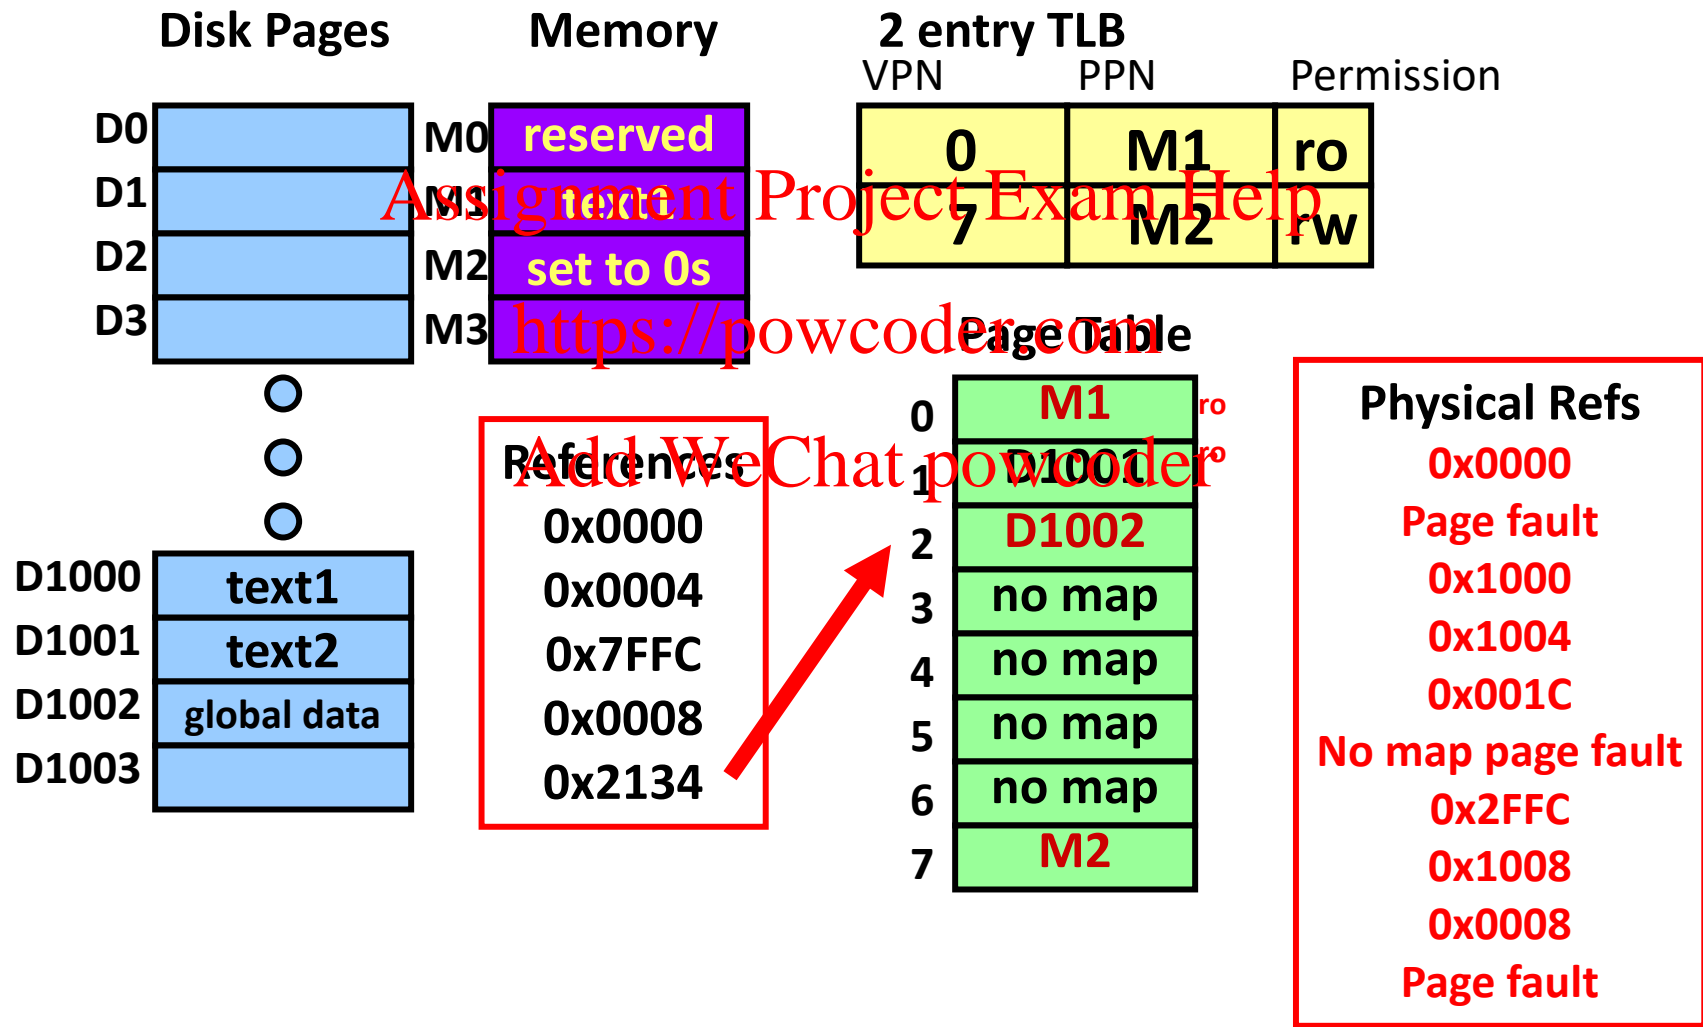
Add WeChat powcoder



Assignment Project Exam Help

Reference 2134

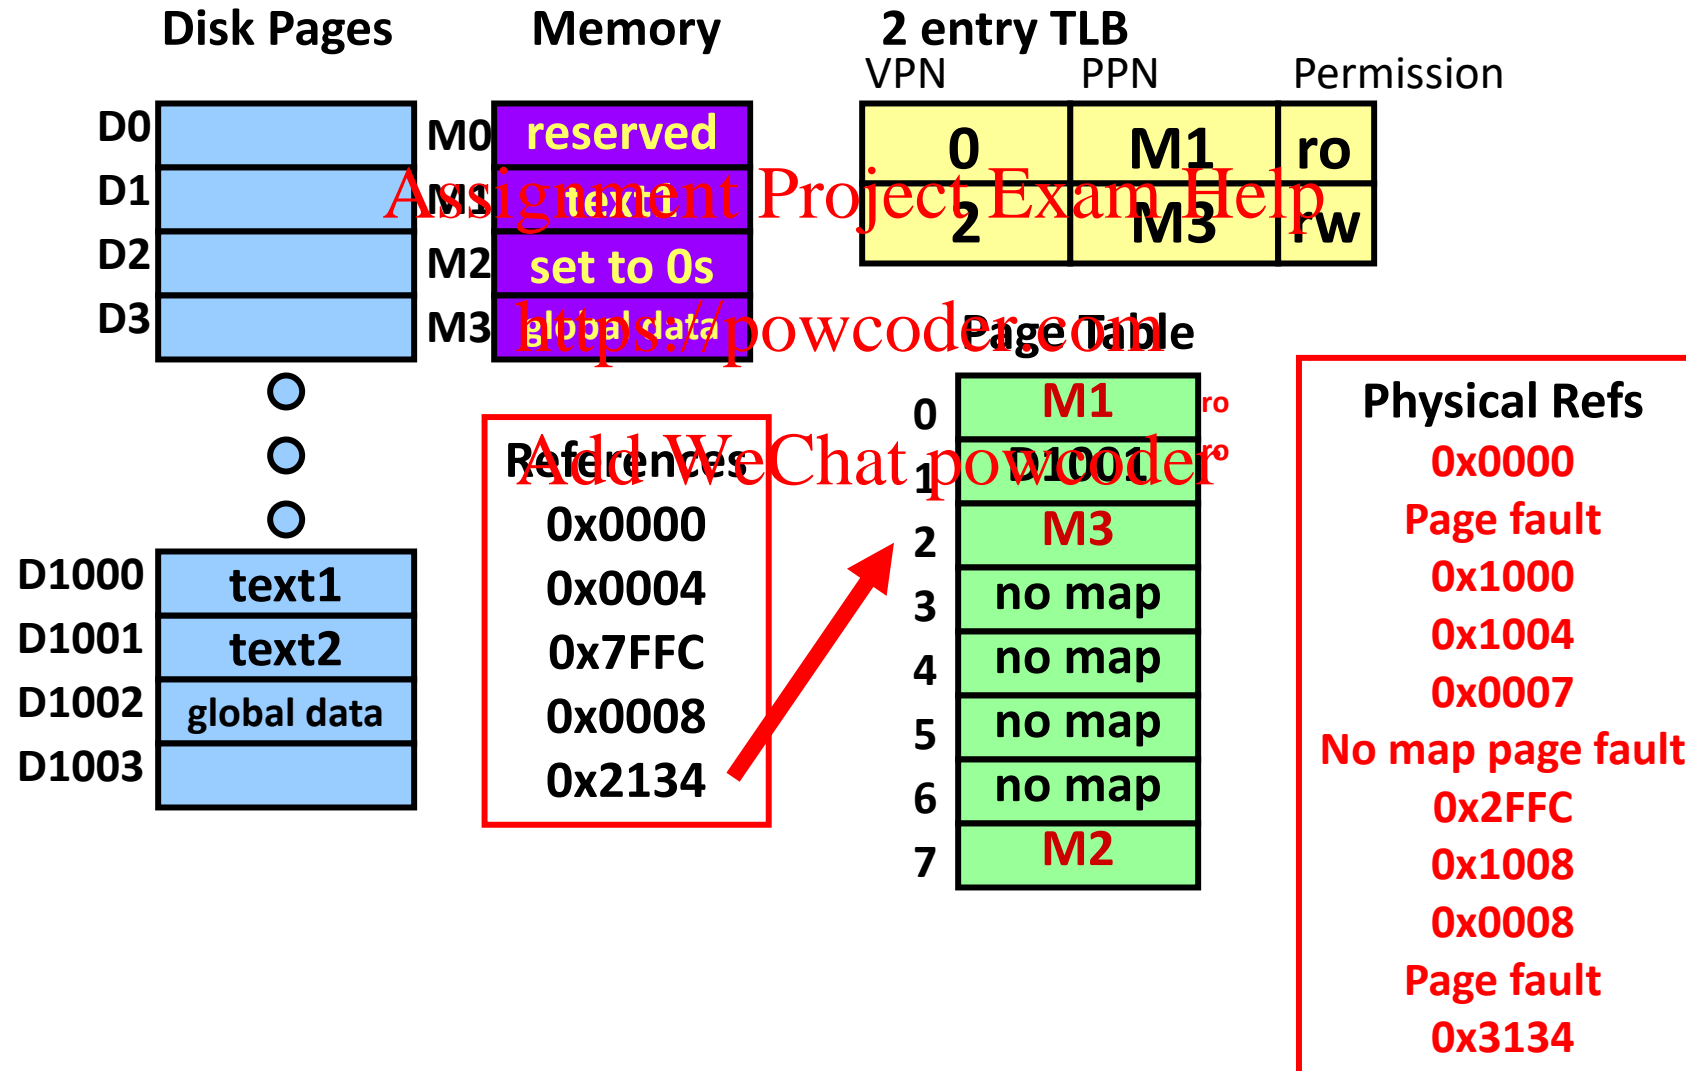
Add WeChat powcoder



Assignment Project Exam Help

Reference 2134

Add WeChat powcoder



Assignment Project Exam Help

Add WeChat powcoder

Can we skip address translation?

<https://powcoder.com>

Virtually-addressed caches

Add WeChat powcoder

Assignment Project Exam Help

Address translation in a pipeline

Add WeChat powcoder

Load/store: Memory stage

Instruction fetch: Fetch stage

Assignment Project Exam Help

When to do address translation?

After VA is computed, but before memory access is performed

Add WeChat powcoder

Is it possible to skip address translation for some memory accesses?

Yes. Answer: Virtually-addressed caches.


Assignment Project Exam Help

Option 1: Address translation *before* cache access

Physically-addressed Cache

✗ Slower

✓ Low complexity



Use ***physical*** address to access cache
(tag, set index, and block offset bits)

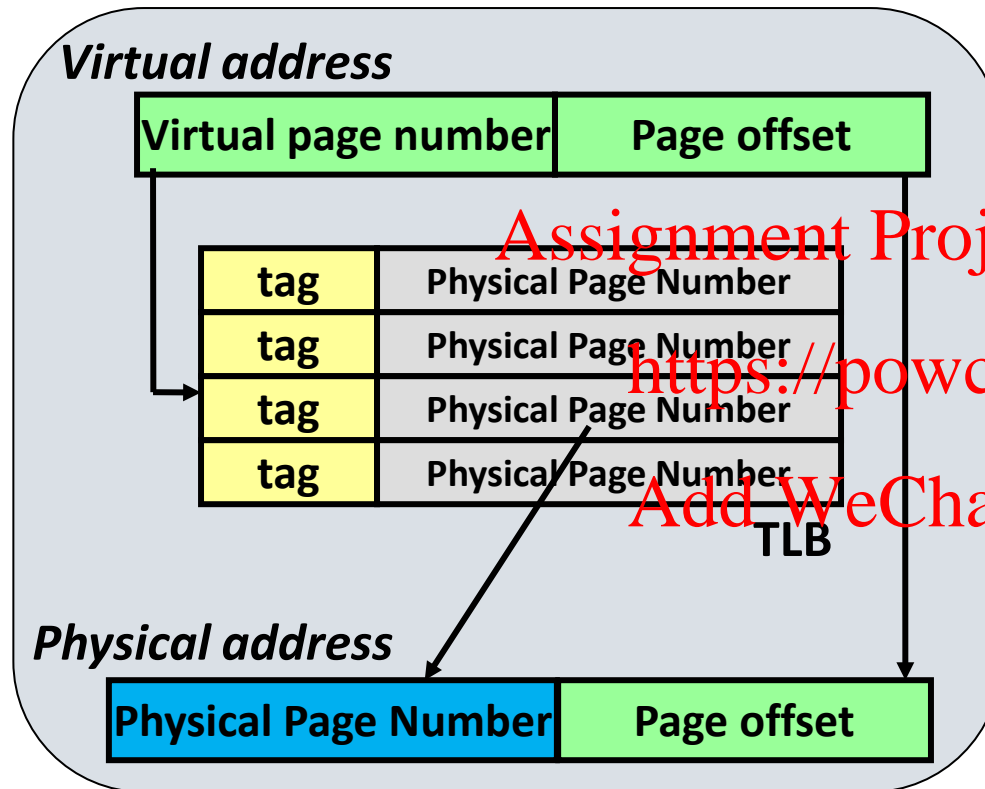
Address translation only for all accesses

Assignment Project Exam Help


Physically addressed caches

Add WeChat powcoder

Step 1: Virtual address to Physical



Step 2: Access the cache with physical address obtained from translation



Assignment Project Exam Help

Option 2: Address translation *after* cache access


Virtually-addressed Cache

✗ High complexity
(process isolation is hard)

✓ Faster

Use **virtual** address to access cache
(tag, set index, and block offset bits)

Address translation only for cache misses



Assignment Project Exam Help

Option 2: Address translation *after* cache access

Virtually-addressed Cache


✗ High complexity
(process isolation is hard)

✓ Faster

TLB can be accessed in parallel
with cache lookup.

Saves time on a cache miss, where
physical address is needed to
access main memory

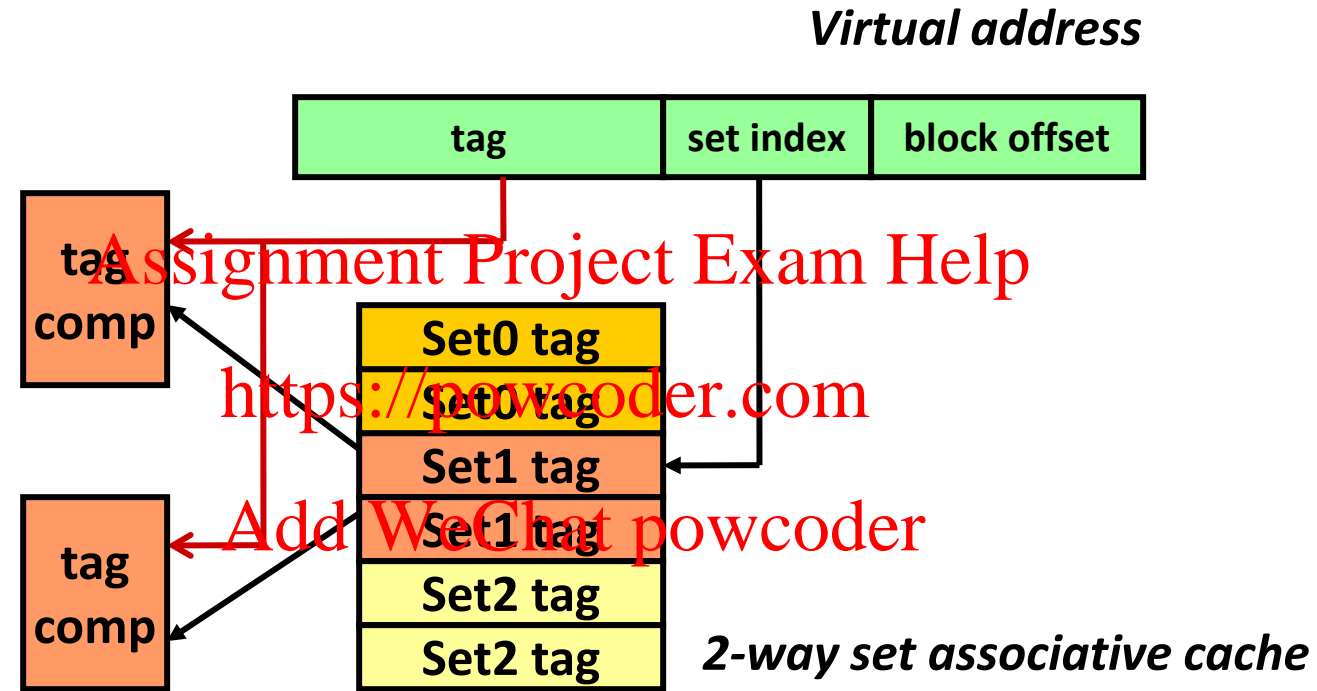
But wasteful TLB accesses on cache hits



Assignment Project Exam Help

Virtually addressed caches

Add WeChat powcoder



Address translation: Before or After Cache Access

Add WeChat powcoder

Physically-addressed cache

Generate virtual address ->

Access TLB ->

Access cache ->

if cache miss, access main memory

Before Cache

Virtually-addressed cache

Generate virtual address ->

Access cache ->

if cache miss, access TLB ->

Access main memory

After Cache

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

Tradeoffs

Add WeChat powcoder

Physically-addressed caches: Slow; Simple

Assignment Project Exam Help

Vs

<https://powcoder.com>

Virtually-addressed caches: Fast; Complex

Add WeChat powcoder

Assignment Project Exam Help

Physical Vs Virtual Caches: Latency

Add WeChat powcoder

Physically-addressed caches

Cache is accessed with physical address (after VM translations).

- Slow. Cache can be accessed only after address translation
- Inefficient, because all accesses need address translation

Virtually-addressed caches

Cache is accessed with virtual address (before VM translation).

- Fast. Skips address translation for cache hits.
- Efficient, because only cache misses need address translation

Assignment Project Exam Help

Physical Vs Virtual Caches: Complexity

Add WeChat powcoder

Problem for virtually-addressed cache:

The same virtual address refers to different physical addresses
in two different processes.

Assignment Project Exam Help

<https://powcoder.com>

To ensure process isolation, on a context switch:

Add WeChat powcoder

Virtual cache need to be invalidated. Dirty cache blocks written back.

Physical cache need not be invalidated.

So, physical cache incurs fewer cache misses if context switches are very frequent
(but generally they are not)

Assignment Project Exam Help

Add WeChat powcoder

Assignment Project Exam Help

Typically, in modern processors

Level-1 (L1) caches are virtually-addressed. L1 needs to be fast.

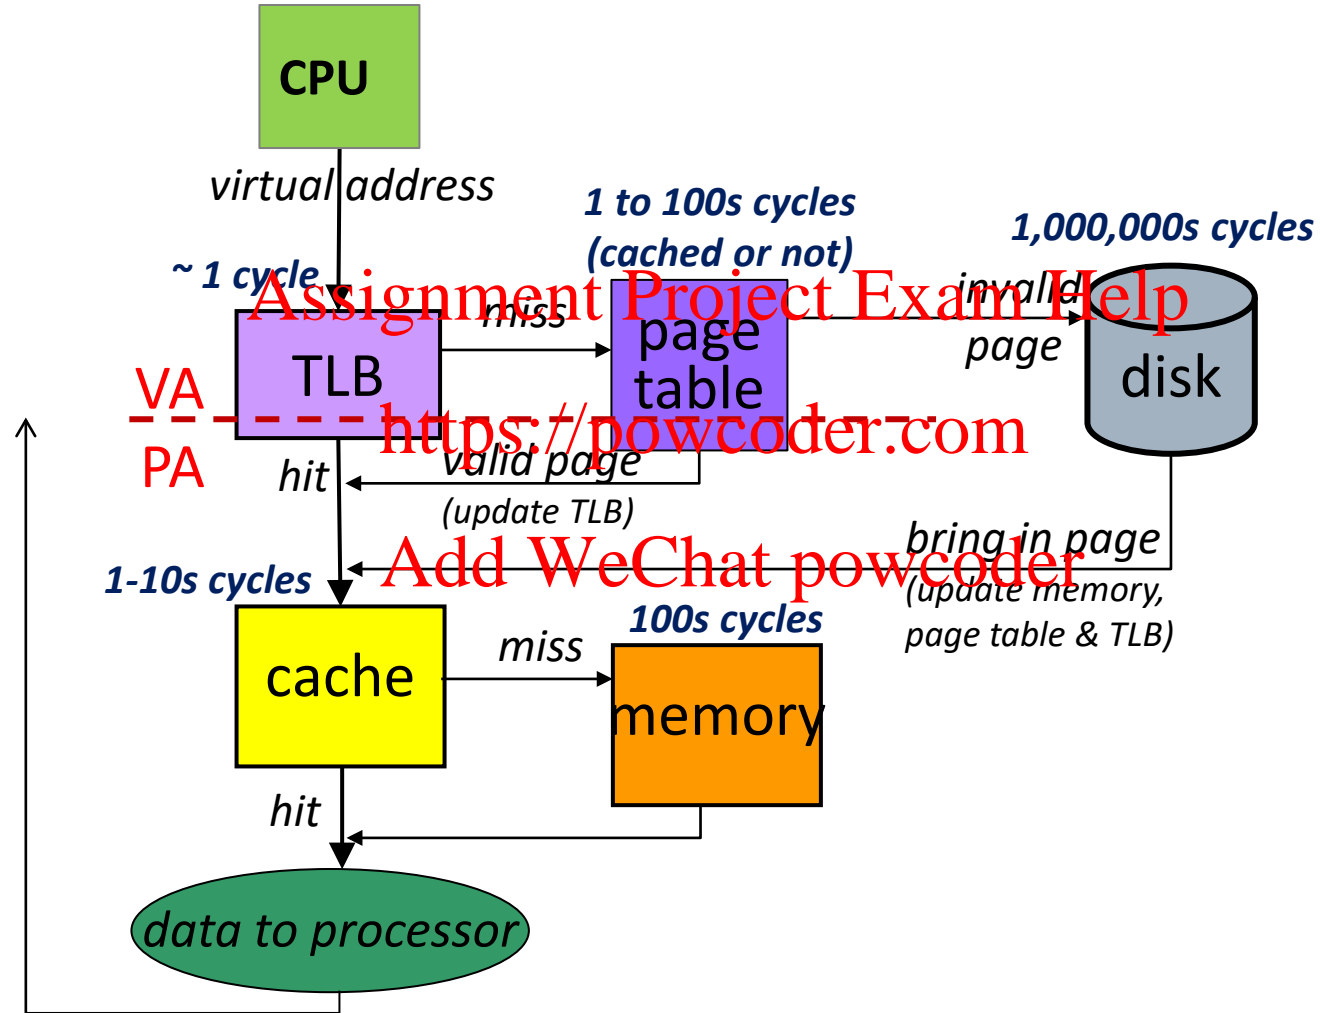
L2 and L3 are physically-addressed. Larger structures.

Checking TLB before accessing them does not significantly affect their latency.

Assignment Project Exam Help

Physically addressed caches: detailed flow

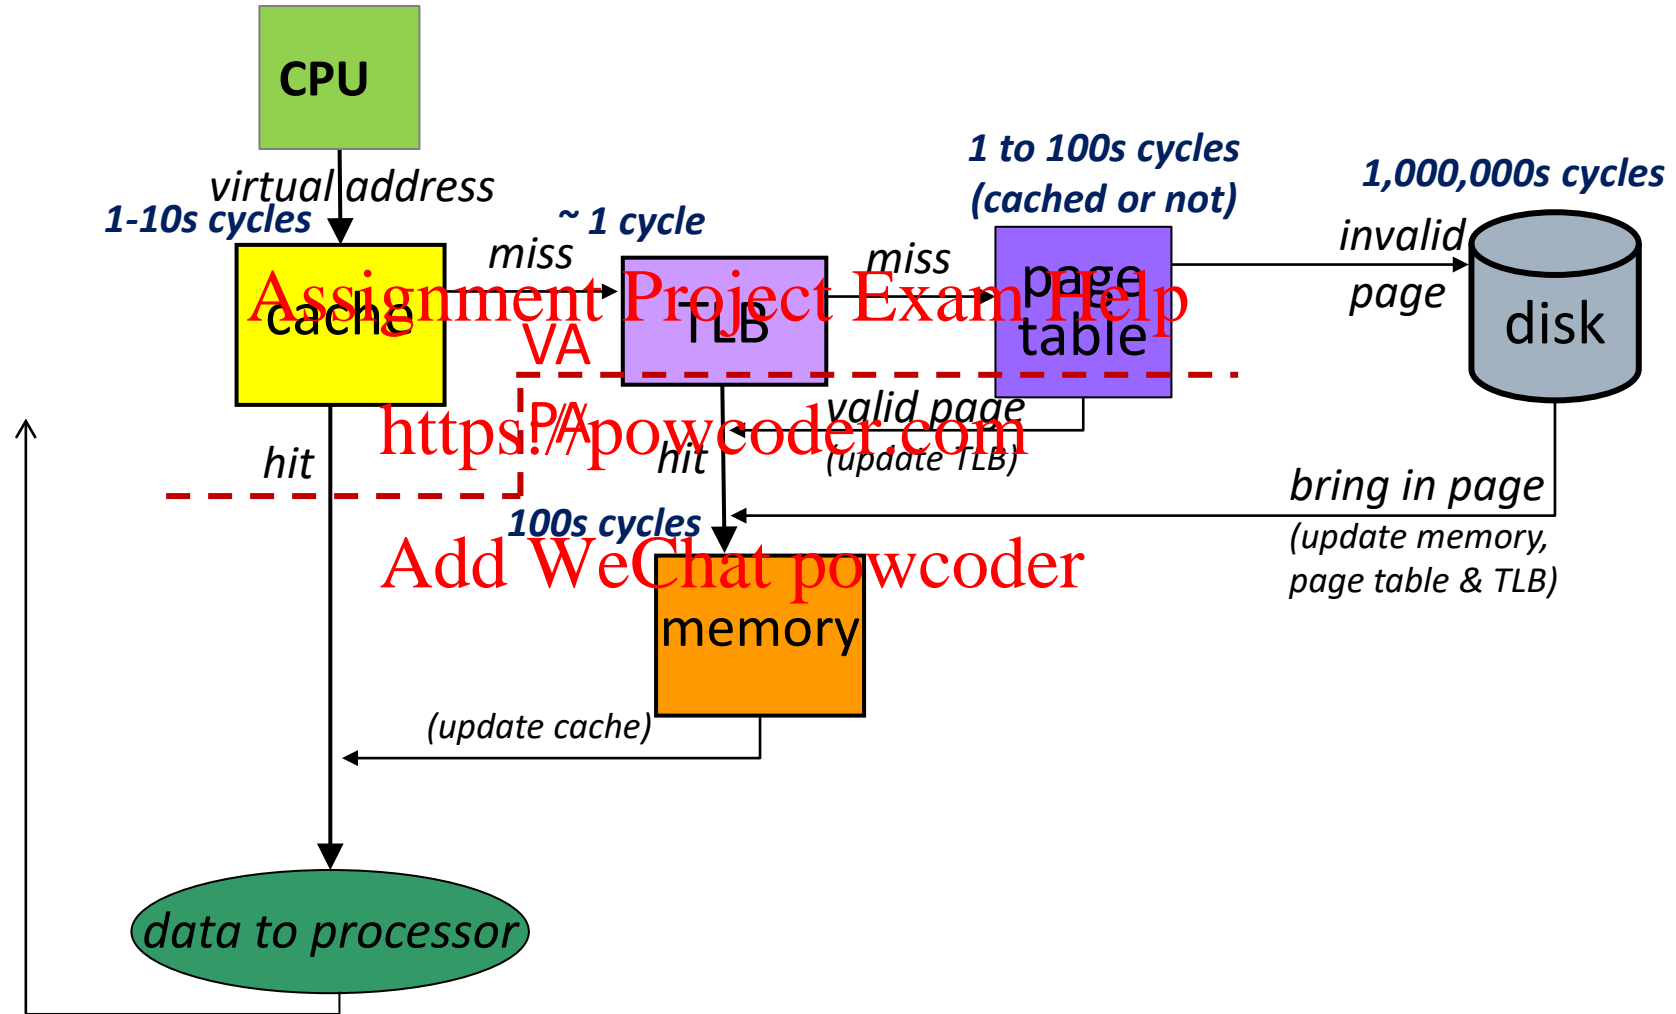
Add WeChat powcoder



Assignment Project Exam Help

Virtually addressed caches: detailed flow

Add WeChat powcoder



Assignment Project Exam Help

OS Support for Virtual Memory

Add WeChat powcoder

OS must be able to modify the page table register, update page table values, etc.

Assignment Project Exam Help
To enable the OS to do this, **BUT** not the user program, we have different execution modes for a process.

- **Executive** (or supervisor or kernel level) permissions and
- **User level** permissions.

Assignment Project Exam Help

References (not part of Course Syllabus)

Add WeChat powcoder

See how Intel's memory management hardware works, Intel x86 Software Manual:
<http://www.intel.com/content/www/us/en/architecture-and-technology/64-ia-32-architectures-software-developer-vol-3a-part-1-manual.html>

Chapter 4 is on Paging

Assignment Project Exam Help

Linux page table management:

<https://www.kernel.org/doc/gupman.html> <https://powcoder.com/understand006.html>

Add WeChat powcoder